

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
sns.set_theme(color_codes=True)
pd.set_option('display.max_columns', None)
```

```
In [2]: df = pd.read_csv('supply_chain_data.csv')
df.head()
```

Out[2]:

	Product type	SKU	Price	Availability	Number of products sold	Revenue generated	Customer demographics	Stock levels	Lead times	Order quantities	Shipping times	Shipping costs
0	haircare	SKU0	69.808006	55	802	8661.996792	Non-binary	58	7	96	4	Ca
1	skincare	SKU1	14.843523	95	736	7460.900065	Female	53	30	37	2	Ca
2	haircare	SKU2	11.319683	34	8	9577.749626	Unknown	1	10	88	2	Ca
3	skincare	SKU3	61.163343	68	83	7766.836426	Non-binary	23	13	59	6	Ca
4	skincare	SKU4	4.805496	26	871	2686.505152	Non-binary	5	3	56	8	Ca

Data Preprocessing Part 1

```
In [3]: df.dtypes
```

```
Out[3]: Product type      object
SKU                      object
Price                   float64
Availability             int64
Number of products sold int64
Revenue generated       float64
Customer demographics   object
Stock levels            int64
Lead times              int64
Order quantities        int64
Shipping times          int64
Shipping carriers       object
Shipping costs          float64
Supplier name           object
Location                object
Lead time               int64
Production volumes      int64
Manufacturing lead time int64
Manufacturing costs     float64
Inspection results      object
Defect rates            float64
Transportation modes    object
Routes                  object
Costs                   float64
dtype: object
```

```
In [4]: df.select_dtypes(include='object').nunique()
```

```
Out[4]: Product type      3
        SKU              100
        Customer demographics      4
        Shipping carriers      3
        Supplier name      5
        Location      5
        Inspection results      3
        Transportation modes      4
        Routes      3
        dtype: int64
```

```
In [5]: df.shape
```

```
Out[5]: (100, 24)
```

```
In [6]: #Drop SKU Column because this is just supply chain id
        df.drop(columns=['SKU'], inplace=True)
        df.shape
```

```
Out[6]: (100, 23)
```

```
In [7]: df.nunique()
```

```
Out[7]: Product type      3
        Price      100
        Availability      63
        Number of products sold      96
        Revenue generated      100
        Customer demographics      4
        Stock levels      65
        Lead times      29
        Order quantities      61
        Shipping times      10
        Shipping carriers      3
        Shipping costs      100
        Supplier name      5
        Location      5
        Lead time      29
        Production volumes      96
        Manufacturing lead time      30
        Manufacturing costs      100
        Inspection results      3
        Defect rates      100
        Transportation modes      4
        Routes      3
        Costs      100
        dtype: int64
```

Exploratory Data Analysis


```
In [8]: # List of categorical variables to plot
cat_vars = ['Product type', 'Customer demographics', 'Shipping carriers', 'Supplier name', 'Location',
            'Inspection results', 'Transportation modes', 'Routes']

# create figure with subplots
fig, axs = plt.subplots(nrows=2, ncols=4, figsize=(20, 10))
axs = axs.flatten()

# create barplot for each categorical variable
for i, var in enumerate(cat_vars):
    sns.barplot(x=var, y='Costs', data=df, ax=axs[i], estimator=np.mean)
    axs[i].set_xticklabels(axs[i].get_xticklabels(), rotation=90)

# adjust spacing between subplots
fig.tight_layout()

# show plot
plt.show()
```



```

In [9]: num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
                    'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
                    'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
                    'Defect rates']

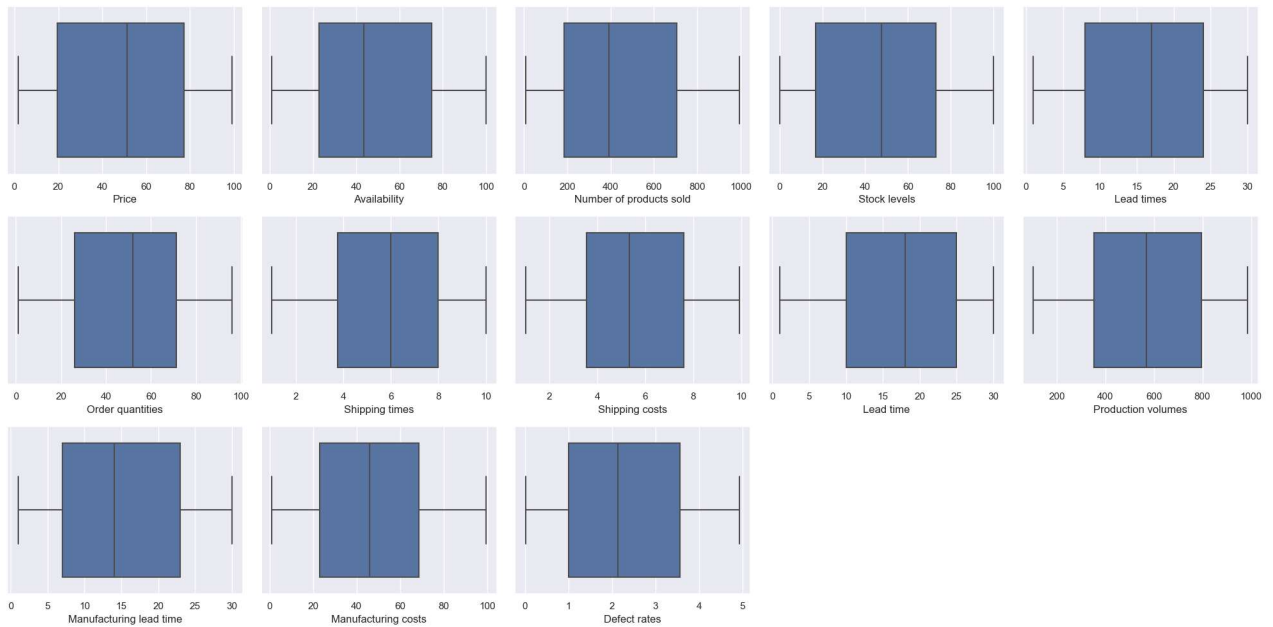
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.boxplot(x=var, data=df, ax=axs[i])

# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])

fig.tight_layout()
plt.show()

```



```

In [10]: num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
                    'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
                    'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
                    'Defect rates']

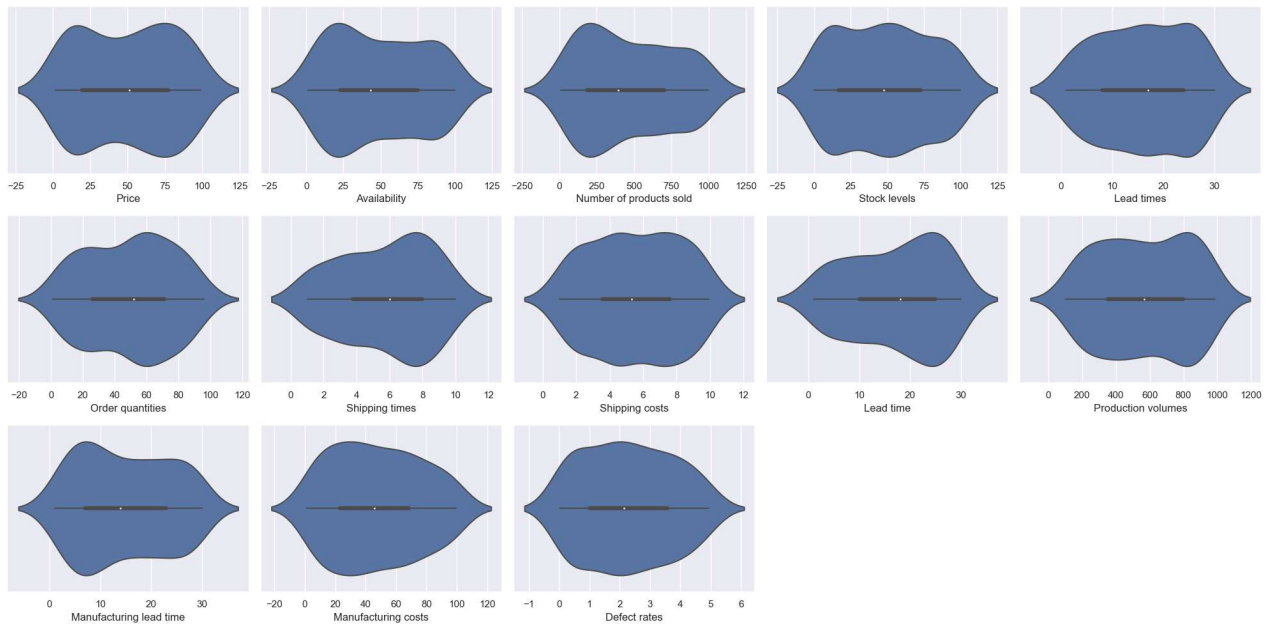
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.violinplot(x=var, data=df, ax=axs[i])

# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])

fig.tight_layout()
plt.show()

```



```

In [11]: num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
                    'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
                    'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
                    'Defect rates']

fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()

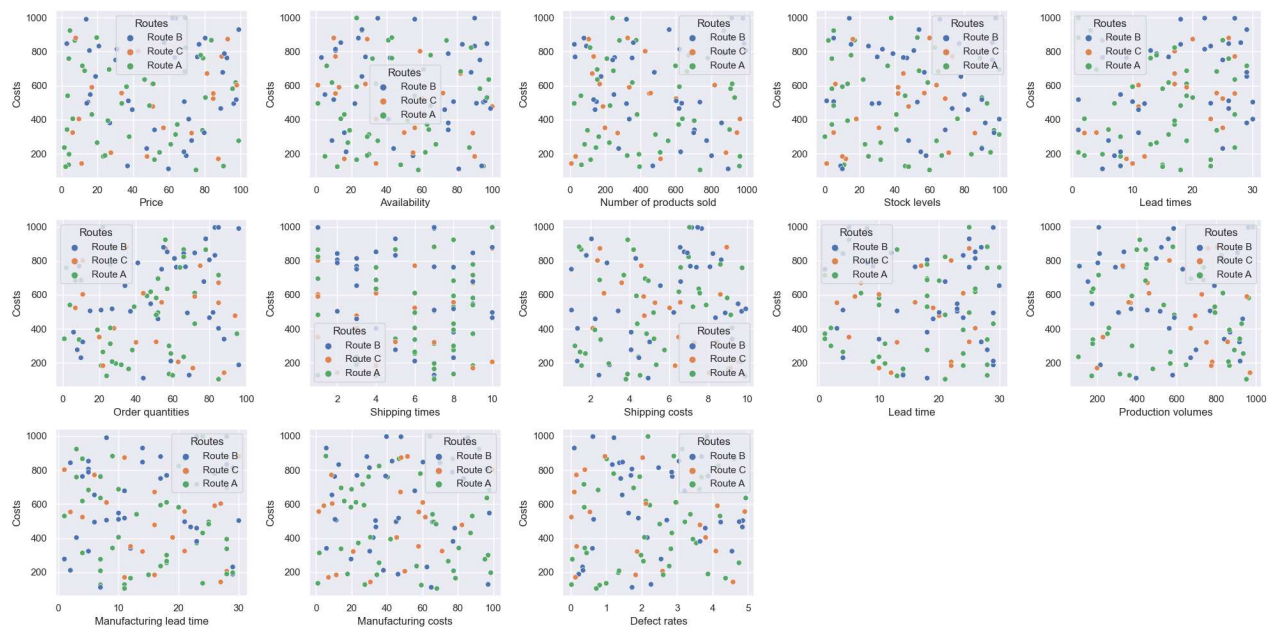
for i, var in enumerate(num_vars):
    sns.scatterplot(x=var, y='Costs', hue='Routes', data=df, ax=axs[i])

# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])

fig.tight_layout()

plt.show()

```



```

In [12]: num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
                    'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
                    'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
                    'Defect rates']

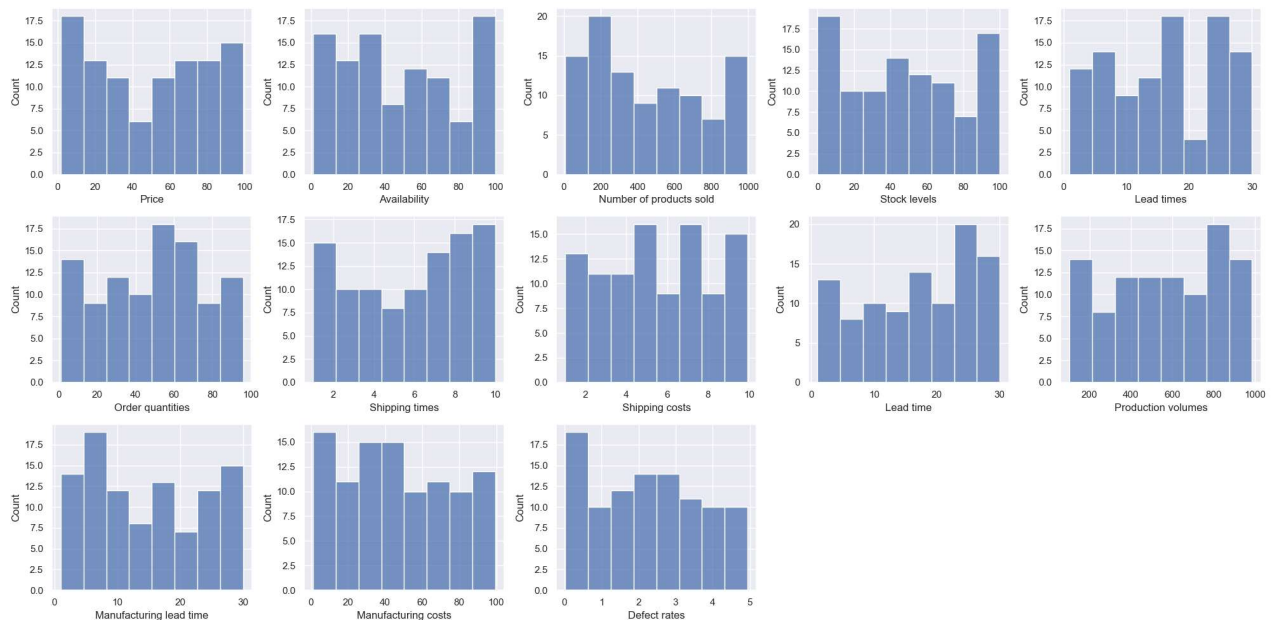
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.histplot(x=var, data=df, ax=axs[i])

# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])

fig.tight_layout()
plt.show()

```



Data Preprocessing Part 2

```

In [13]: #Check the missing value
check_missing = df.isnull().sum() * 100 / df.shape[0]
check_missing[check_missing > 0].sort_values(ascending=False)

```

```

Out[13]: Series([], dtype: float64)

```

Label Encoding for Object datatypes

```
In [14]: # Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:

    # Print the column name and the unique values
    print(f"{col}: {df[col].unique()}")
```

Product type: ['haircare' 'skincare' 'cosmetics']
Customer demographics: ['Non-binary' 'Female' 'Unknown' 'Male']
Shipping carriers: ['Carrier B' 'Carrier A' 'Carrier C']
Supplier name: ['Supplier 3' 'Supplier 1' 'Supplier 5' 'Supplier 4' 'Supplier 2']
Location: ['Mumbai' 'Kolkata' 'Delhi' 'Bangalore' 'Chennai']
Inspection results: ['Pending' 'Fail' 'Pass']
Transportation modes: ['Road' 'Air' 'Rail' 'Sea']
Routes: ['Route B' 'Route C' 'Route A']

```
In [15]: from sklearn import preprocessing

# Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:

    # Initialize a LabelEncoder object
    label_encoder = preprocessing.LabelEncoder()

    # Fit the encoder to the unique values in the column
    label_encoder.fit(df[col].unique())

    # Transform the column using the encoder
    df[col] = label_encoder.transform(df[col])

    # Print the column name and the unique encoded values
    print(f"{col}: {df[col].unique()}")
```

Product type: [1 2 0]
Customer demographics: [2 0 3 1]
Shipping carriers: [1 0 2]
Supplier name: [2 0 4 3 1]
Location: [4 3 2 0 1]
Inspection results: [2 0 1]
Transportation modes: [2 0 1 3]
Routes: [1 2 0]


```
In [16]: df.dtypes
```


```
Out[16]: Product type          int32
Price                float64
Availability          int64
Number of products sold  int64
Revenue generated      float64
Customer demographics   int32
Stock levels           int64
Lead times             int64
Order quantities        int64
Shipping times          int64
Shipping carriers       int32
Shipping costs          float64
Supplier name           int32
Location               int32
Lead time               int64
Production volumes      int64
Manufacturing lead time int64
Manufacturing costs      float64
Inspection results       int32
Defect rates             float64
Transportation modes     int32
Routes                  int32
Costs                    float64
dtype: object
```

There's no outlier so we dont have to remove it

Correlation Heatmap

```
In [17]: #Correlation Heatmap
plt.figure(figsize=(20, 16))
sns.heatmap(df.corr(), fmt='.2g', annot=True)
```

Out[17]: <AxesSubplot:>



Train test Split

```
In [50]: X = df.drop('Costs', axis=1)
y = df['Costs']
```

```
In [51]: #test size 20% and train size 80%
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,random_state=0)
```

Decision Tree Regressor

```
In [52]: from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import load_boston

# Create a DecisionTreeRegressor object
dtree = DecisionTreeRegressor()

# Define the hyperparameters to tune and their values
param_grid = {
    'max_depth': [2, 4, 6, 8],
    'min_samples_split': [2, 4, 6, 8],
    'min_samples_leaf': [1, 2, 3, 4],
    'max_features': ['auto', 'sqrt', 'log2'],
    'random_state': [0, 7, 42]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(dtree, param_grid, cv=5, scoring='neg_mean_squared_error')

# Fit the GridSearchCV object to the data
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print(grid_search.best_params_)

{'max_depth': 2, 'max_features': 'sqrt', 'min_samples_leaf': 3, 'min_samples_split': 2, 'random_state': 0}
```

```
In [60]: from sklearn.tree import DecisionTreeRegressor
dtree = DecisionTreeRegressor(random_state=0, max_depth=2, max_features='sqrt', min_samples_leaf=
dtree.fit(X_train, y_train)
```

```
Out[60]: DecisionTreeRegressor(max_depth=2, max_features='sqrt', min_samples_leaf=3,
                                random_state=0)
```

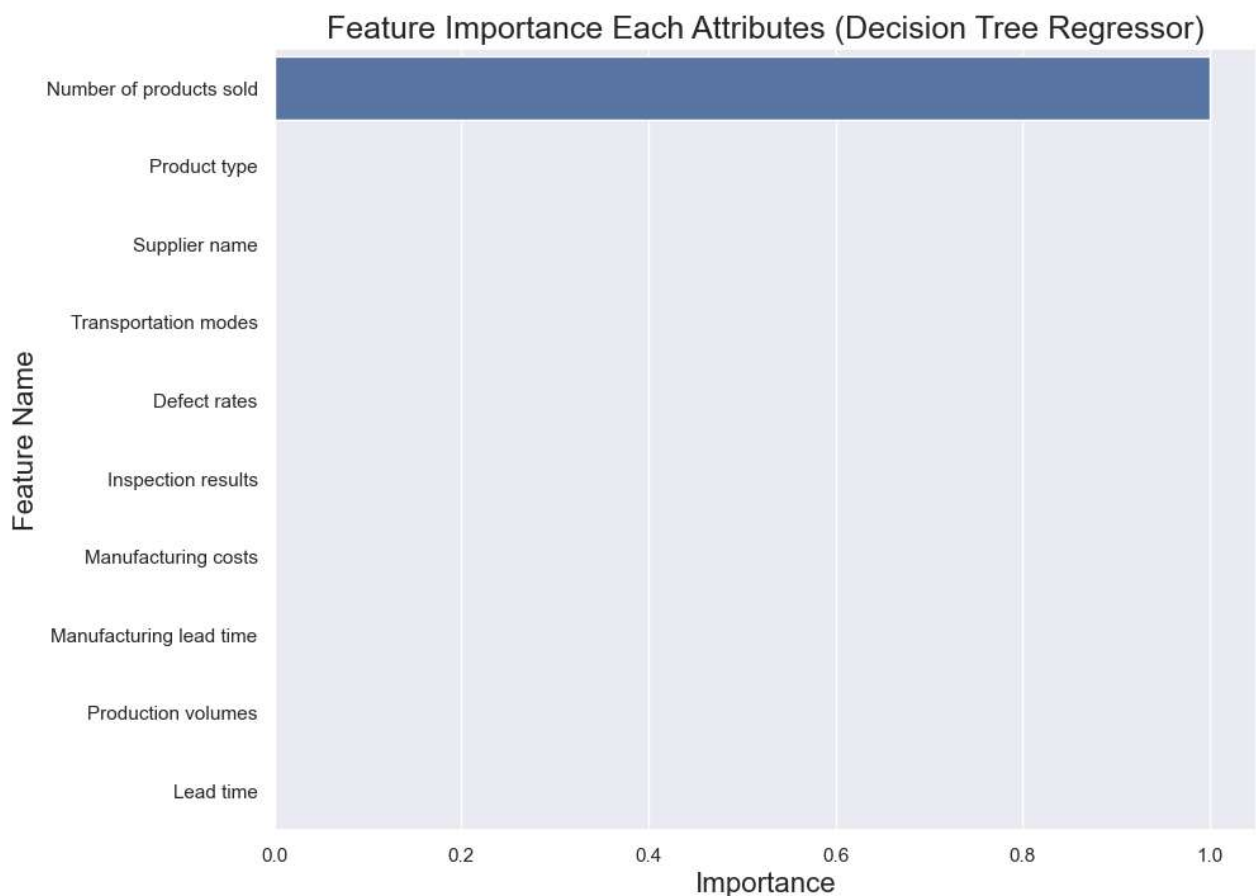
```
In [61]: from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = dtree.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))
```

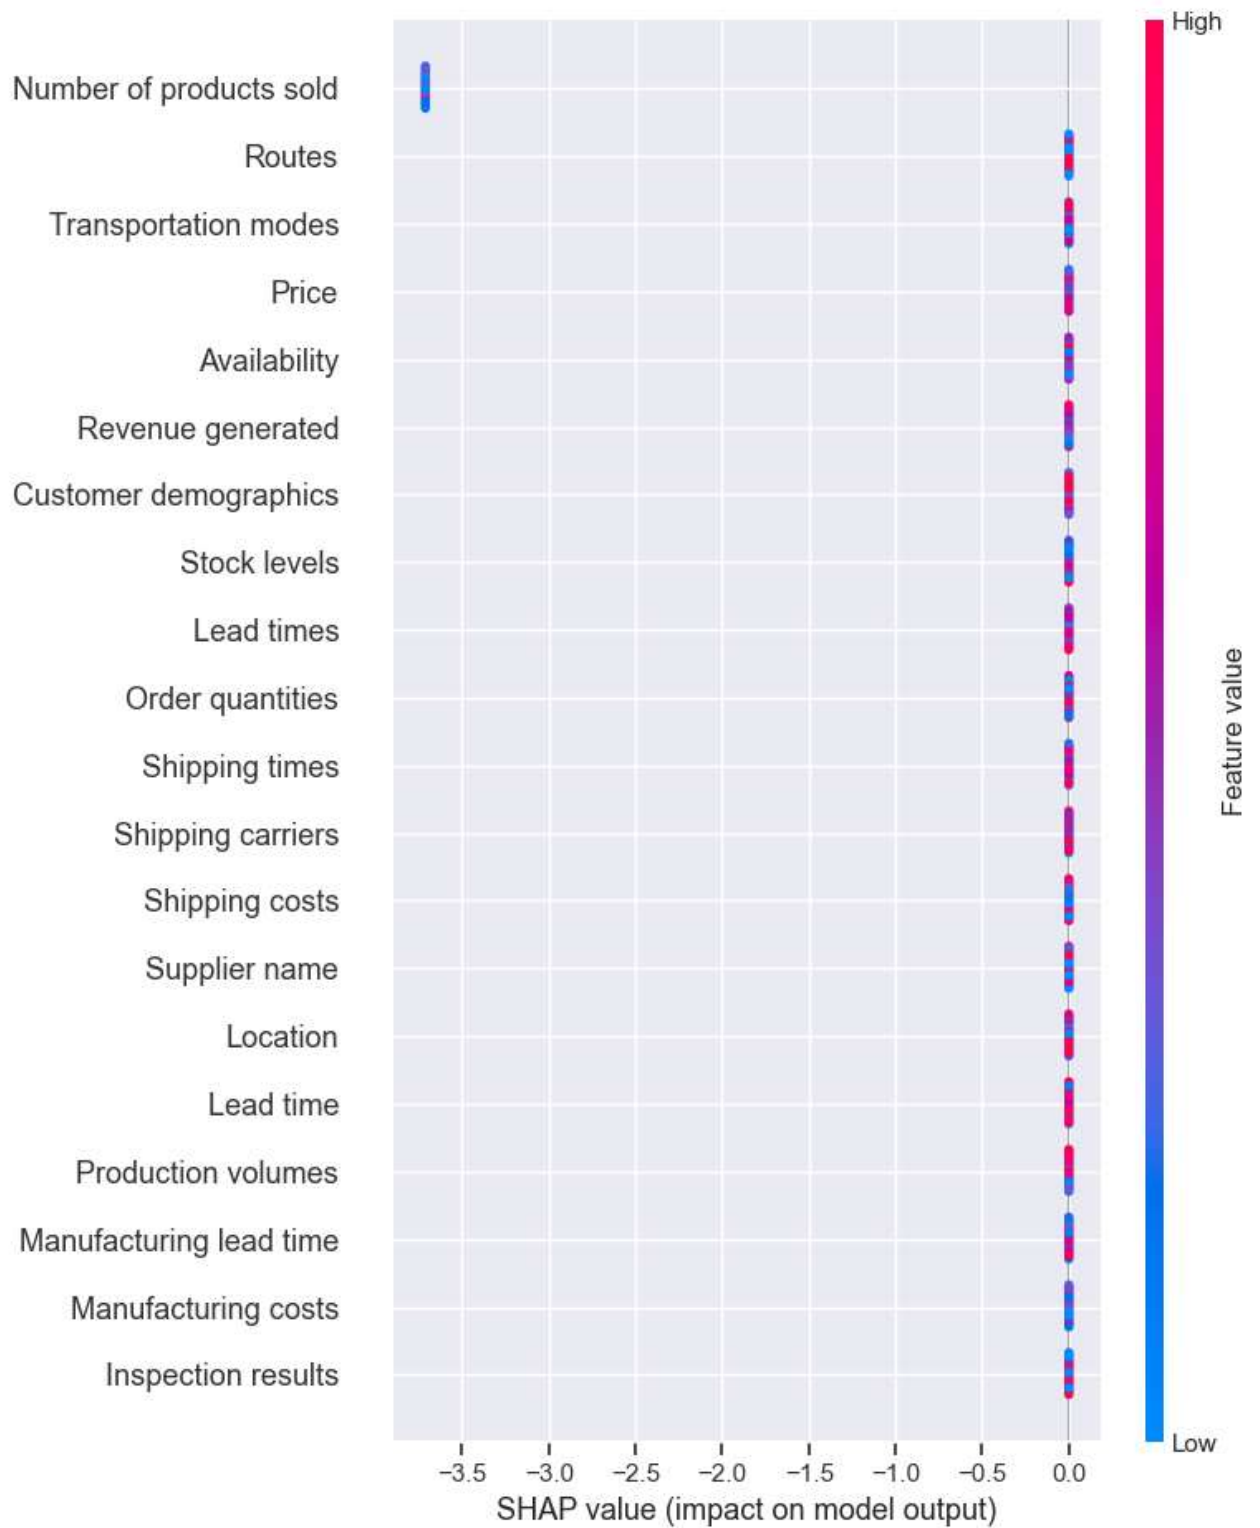
```
MAE is 248.4413893861546
MAPE is 0.5893818876444419
MSE is 72806.47766651674
R2 score is -0.08647889188367719
RMSE score is 269.8267549123266
```

```
In [62]: imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": dtree.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (Decision Tree Regressor)', fontsize=18)
plt.xlabel('Importance', fontsize=16)
plt.ylabel('Feature Name', fontsize=16)
plt.show()
```



```
In [63]: import shap
explainer = shap.TreeExplainer(dtree)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```



Random Forest Regressor

```
In [57]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV

# Create a Random Forest Regressor object
rf = RandomForestRegressor()

# Define the hyperparameter grid
param_grid = {
    'max_depth': [3, 5, 7, 9],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt'],
    'random_state': [0, 7, 42]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='r2')

# Fit the GridSearchCV object to the training data
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print("Best hyperparameters: ", grid_search.best_params_)

Best hyperparameters: {'max_depth': 3, 'max_features': 'sqrt', 'min_samples_leaf': 2, 'min_samples_split': 5, 'random_state': 0}
```

```
In [64]: from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(random_state=0, max_depth=3, min_samples_split=5, min_samples_leaf=2,
                           max_features='sqrt')
rf.fit(X_train, y_train)
```

```
Out[64]: RandomForestRegressor(max_depth=3, max_features='sqrt', min_samples_leaf=2,
                               min_samples_split=5, random_state=0)
```

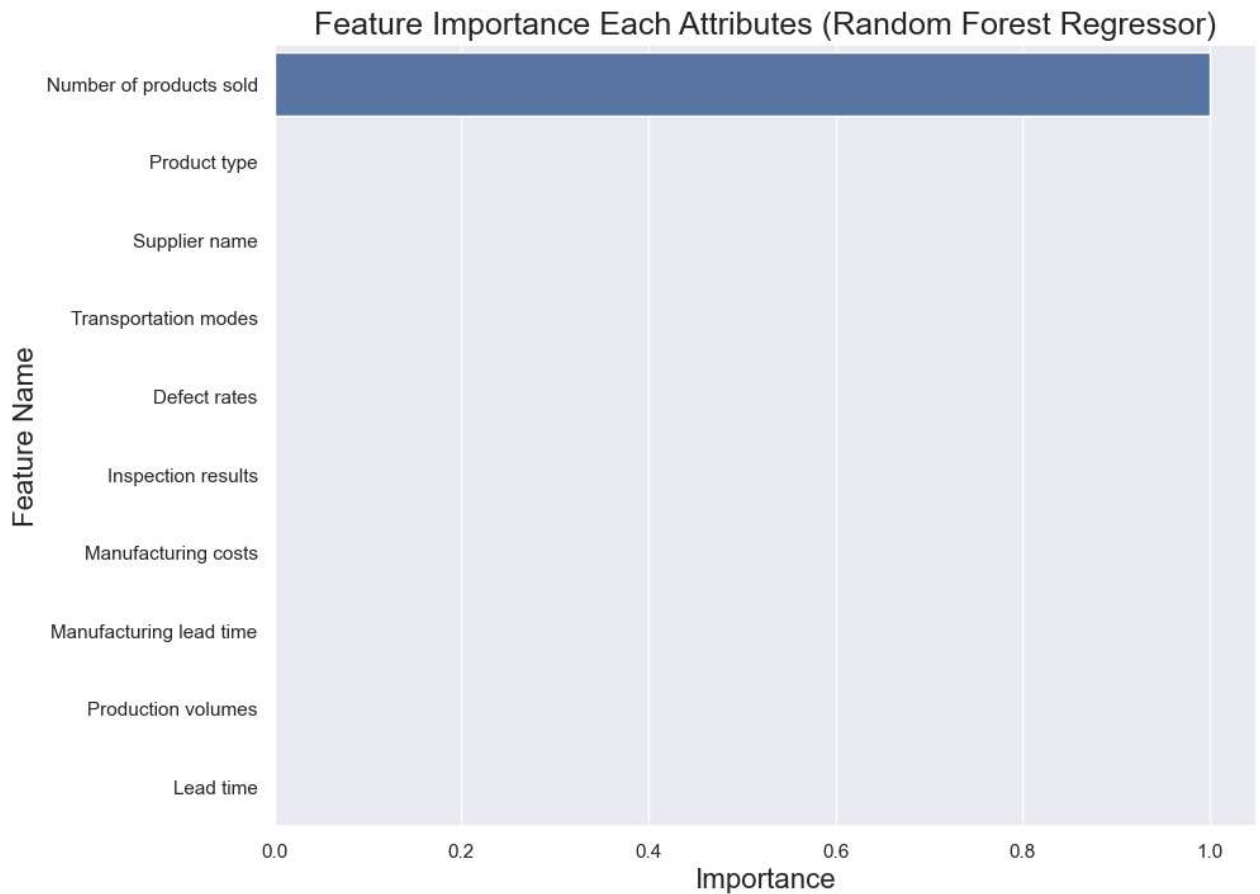
```
In [65]: from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = rf.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))
```

```
MAE is 247.33969719962744
MAPE is 0.6029768224226728
MSE is 71899.28833186119
R2 score is -0.07294105713825938
RMSE score is 268.14042651540103
```

```
In [66]: imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": dtree.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (Random Forest Regressor)', fontsize=18)
plt.xlabel('Importance', fontsize=16)
plt.ylabel('Feature Name', fontsize=16)
plt.show()
```



```
In [67]: import shap
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```



AdaBoost Regressor


```
In [68]: from sklearn.ensemble import AdaBoostRegressor
from sklearn.model_selection import GridSearchCV

# Create an AdaBoost Regressor object
ada = AdaBoostRegressor()

# Define the hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 150, 200],
    'learning_rate': [0.01, 0.1, 1],
    'loss': ['linear', 'square', 'exponential'],
    'random_state': [0, 7, 42]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(ada, param_grid, cv=5, scoring='r2')

# Fit the GridSearchCV object to the training data
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print("Best hyperparameters: ", grid_search.best_params_)
```

Best hyperparameters: {'learning_rate': 1, 'loss': 'linear', 'n_estimators': 50, 'random_state': 7}

```
In [69]: from sklearn.ensemble import AdaBoostRegressor
ada = AdaBoostRegressor(random_state=7, n_estimators=50, learning_rate=1, loss='linear')
ada.fit(X_train, y_train)
```

Out[69]: AdaBoostRegressor(learning_rate=1, random_state=7)

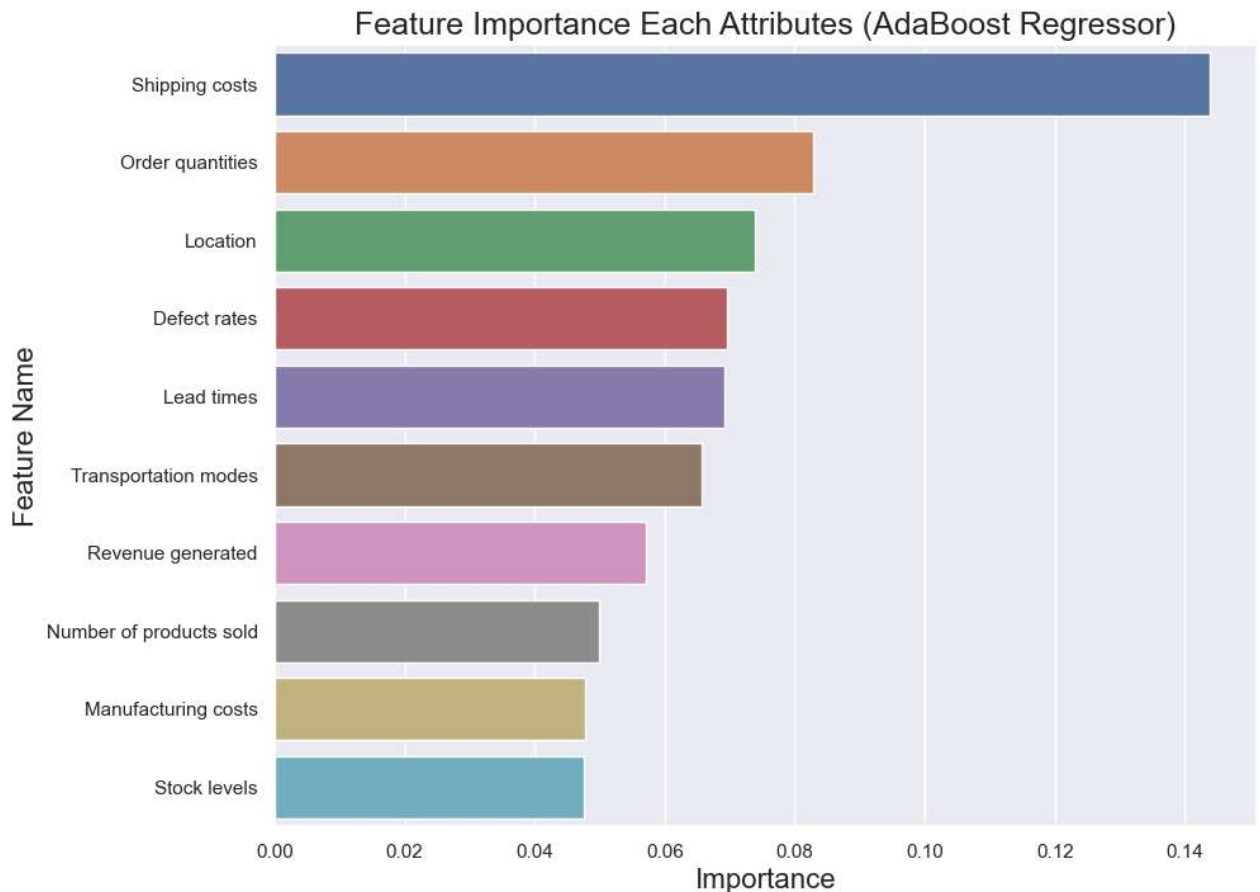
```
In [70]: from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = ada.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))
```

MAE is 255.28612180541396
MAPE is 0.5859844238678936
MSE is 78800.74415400976
R2 score is -0.17593032834538103
RMSE score is 280.71470241868303

```
In [72]: imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": ada.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (AdaBoost Regressor)', fontsize=18)
plt.xlabel('Importance', fontsize=16)
plt.ylabel('Feature Name', fontsize=16)
plt.show()
```



All of the Algorithms got bad R2 Score and MAPE Score even with hyperparameter tuning because we only have 100 data and the distribution is spread